

Computer Lab 2 - Bayesian Nonparametrics

The labs are the only examination, so you should do the labs **individually**.

You can use any programming language you prefer, but do **submit the code**.

Submit a readable report in **PDF** (no Word documents!) or a **JuPyteR notebook**

1. **Dirichlet process.** Let y_i be iid observations from some unknown probability distribution P . Let $P \sim DP(\alpha P_0)$ a priori, where P_0 is the $N(20, 3^2)$ base measure.
 - (a) Plot the empirical CDF (cumulative distribution function) for the Galaxy velocity data (as given in the file `GalaxyData.dat`).
 - (b) Use the stick-breaking construction to simulate draws of P from the *prior* of $P \sim DP(\alpha P_0)$. Plot the CDF for some of the draws in a graph to illustrate the prior variation. Plot also the prior mean of P . Try this for $\alpha = 1$, $\alpha = 10$ and $\alpha = 100$.
 - (c) Use the stick-breaking construction to simulate draws of P from the *posterior* of P based on the Galaxy data. Plot the CDF for some of the posterior draws in a graph to illustrate the posterior variation. Plot also the posterior mean of P . Again, try this for $\alpha = 1$, $\alpha = 10$ and $\alpha = 100$. Compare with the prior.
2. **Dirichlet process mixture.** Consider again the Galaxy velocity data, but now modelled by an infinite DP mixture model with a normal/Gaussian density kernel

$$\mathcal{K}(y_i|\theta_i) = \frac{1}{\sqrt{2\pi\sigma_i^2}} \exp\left(-\frac{1}{2\sigma_i^2}(y_i - \mu_i)^2\right)$$

where $\theta_i = (\mu_i, \sigma_i^2)$ are the parameters for the i th observation y_i . Let $(\mu_i, \sigma_i^2) \sim P$ where $P \sim DP(\alpha P_0)$ and P_0 is the conjugate prior

$$\mu_i|\sigma_i^2 \sim N(\mu_0, \sigma_i^2/\kappa_0) \text{ and } \sigma_i^2 \sim \text{Inv} - \chi^2(\nu_0, \sigma_0^2).$$

Set $\mu_0 = 20$, $\sigma_0^2 = 20^2$, $\kappa_0 = 5$ and $\nu_0 = 3$.

- (a) Implement the Blocked Gibbs sampler on page 552 in BDA3. [Hint: the results from Page 67-69 from the BDA3 book will be useful].
- (b) Analyze the Galaxy data using the Blocked Gibbs sampler in 2a) above. Plot the posterior distribution of the number of non-empty components. Plot a regular (non-Bayesian) histogram of the data and overlay the fitted DPM density. Investigate the effect of α by performing the analysis with $\alpha = 1$, $\alpha = 10$ and $\alpha = 100$.

- (c) [Bonus, if you have the energy.] Now treat α as unknown with prior $\alpha \sim \text{Gamma}(a_\alpha, b_\alpha)$. Add an updating step for α in the Blocked Gibbs sampler (see page 553 in BDA3). Analyze the Galaxy data again using some suitable values for a_α and b_α . Plot the prior and posterior of α . Note that BDA3 uses the so called *rate parametrization* of the Gamma density where if $X \sim \Gamma(a_\alpha, b_\alpha)$ then the pdf is of the form

$$p(x) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}.$$

Good luck! All problems have solutions almost surely!